

Neural Hangman Solver

Ashutosh Rabia, Indian Institute of Technology Kanpur

1st June 2025

Introduction

Hangman is a classic word-guessing game where a hidden word is represented by blanks, and the player proposes letters one at a time. Correct guesses reveal letters, incorrect guesses accumulate until six mistakes result in a loss. Effective solvers leverage linguistic patterns, letter frequencies, and positional cues to maximize success.

This report presents a *hybrid neural-statistical* Hangman solver that:

- Encodes each partially revealed word (with blanks and boundary markers) into a fixed-length tensor, along with embeddings of previously missed letters.
- Trains a bi-directional LSTM (hidden size 512, dropout) to predict, for each blank position, the most likely letter.
- Computes a positional-bias tensor from a 250 000-word dictionary to favor letters historically common at specific positions.
- Combines a fast “frequency filter” (counting how often each letter appears among dictionary words matching the current pattern) with the LSTM’s predicted probabilities. These scores are linearly weighted (0.6 for neural, 0.4 for frequency) and masked to exclude already-guessed letters.

Intuition

1. Vowel-First Heuristic

Early in each game, if no vowel has appeared and no vowel has yet been guessed, the solver forces a vowel guess in the order {e, a, o, i, u}. English words nearly always contain at least one vowel, forcing a vowel early yields high information gain.

2. Regex-Based Candidate Filtering

Once a vowel is placed (or forced), the solver builds a regex pattern $\text{^p}_1 \text{p}_2 \dots \text{p}_n \text{\$}$, where each p_i is a revealed letter or “.” for a blank. For example, the pattern p_t becomes $\text{^p.t\$}$. All dictionary words of length n matching this pattern form the candidate set. For each unguessed letter c , compute

$$\text{freq_score}(c) = \frac{|\{w \in \text{candidates} : c \in w\}|}{\max_{c'} |\{w \in \text{candidates} : c' \in w\}|} \in [0, 1].$$

If there are no candidates or no unguessed letters, all frequency scores are set to 0.

3. Neural Prediction (Bi-LSTM with Positional Bias)

The bi-directional LSTM views the framed pattern $\{\dots |$ plus an embedding of missed letters. Trained on millions of random “masked + partial-reveal” examples, it learns letter

co-occurrence and positional patterns (e.g., “q” likely followed by “u,” or “th” common at word starts). At inference, the LSTM outputs logits $\ell_{t,c}$ for each time step t and each letter c . For blank positions B , define

$$\text{neural_score}(c) = \sum_{t \in B} \text{Softmax}(\ell_{t,c}),$$

with already-guessed letters zeroed out. A precomputed positional-bias tensor $\text{pos_bias}[c, p]$ ($26 \times \text{embedding_len}$) guides training by adding $\lambda_{\text{pos}} \log(\text{pos_bias}[c, p])$ to the loss when predicting letter c at position p .

4. Score Combination

For each candidate letter c , compute

$$\text{combined_score}(c) = \alpha \text{neural_score}(c) + (1 - \alpha) \text{freq_score}(c), \quad \alpha = 0.6.$$

The solver picks $\arg \max_c \text{combined_score}(c)$. This hybrid approach leverages both neural positional knowledge and real-time frequency statistics.

Methodology

Preprocessing

1. **Load & Split Dictionary:** Read 250 000 words from `words_250000_train.txt`, shuffle randomly, split 95 % (237 500 words) for training, 5 % (12 500 words) for validation.
2. **Frame Words:** For each word w , define `padded` = $\{w \mid$. The start marker `{` and end marker `|` explicitly delimit the word.
3. **Compute Positional Bias:**

- $\text{embedding_len} = (\max_w |w|) + 2$. $\max |w| = 29$, so $\text{embedding_len} = 31$.
- Initialize `counts` $\in R^{26 \times 31}$ to zeros. For each framed training word $\{w \mid$, let $\text{start} = 31 - |\{w \mid|$. For index $i = 0, \dots, |\{w \mid| - 1$, let position $p = \text{start} + i$. If the character at that index is a letter $c \in \{a, \dots, z\}$, increment `counts`[$c - a, p$].
- Normalize columns:

$$\text{pos_bias}[c, p] = \frac{\text{counts}[c, p]}{\sum_{c'=0}^{25} \text{counts}[c', p] + 1e-9}, \quad c \in \{0, \dots, 25\}, p \in \{0, \dots, 30\}.$$

Store as a `torch.FloatTensor(26 × 31)` on GPU.

4. **Generate Synthetic Training Samples:** For each training word w :

- (a) Let `padded` = $\{w \mid$. Choose a reveal-percentage $r \in \{0, 0.2, 0.5, 0.8\}$. Compute $\text{num_reveals} = \lfloor r \cdot |w| \rfloor$. Randomly sample that many interior positions to reveal; add those letters to `guessed_letters`.
- (b) Randomly mask $m \in [1, \max(1, |w| - \text{num_reveals})]$ positions (not already revealed) by replacing them with “.”
- (c) Call

$(\vec{missed_vec}, \text{revealed_mask}, \text{mask_pos}) = \text{encode_input}(\text{masked_string}, w, \text{guessed_letters}, 31)$.

$\vec{\in} \{0, \dots, 29\}^{31}$, $\text{missed_vec} \in \{0, 1\}^{26}$, $\text{revealed_mask} \in \{0, 1\}^{31}$, $\text{mask_pos} \subset \{0, \dots, 30\}$.

- (d) If `mask_pos = ∅`, set `target_pos = 0`, `y_letter = 0`. Otherwise pick `target_pos` randomly from `mask_pos`, find the true letter at that index in `{w |, define y_letter = ord(target_letter) - 97.`
- (e) Return the tuple of tensors:

`X = LongTensor()`, `missed_vec = LongTensor(missed_vec)`, `revealed_mask = LongTensor(revealed_mask)`

These examples populate a custom `HangmanDataset` used by `DataLoader(batch_size=256)`.

Model Architecture & Training

1. Character & Missed-Letters Embedding

- `char_embed = nn.Embedding(30, 64)` maps indices `{0, ..., 29}` (padding, `{, |, “-”`, letters) to 64-dim vectors.
- `missed_embed = nn.Linear(26, 16)` embeds the 26-bit “missed letters” mask into 16 dimensions.

2. Bi-Directional LSTM

- Input size per time step: `64 + 16 = 80`.
- Hidden size per direction: 512. Layers: 2. `self.lstm = nn.LSTM(80, 512, num_layers = 2, batch_first = True, bidirectional = True)`.
- Output per time `t`: (batch, 1024) (512 forward + 512 backward).

3. **Dropout Layer** `self.dropout = nn.Dropout(p = 0.5)`. Applied to LSTM outputs during training.
4. **Fully-Connected Layer** `self.fc = nn.Linear(512 × 2, 26)`. Maps each (1024) vector at time `t` to 26 logits.
5. **Loss with Positional Bias** For a batch of size `B`:

$$\text{logits} = \text{model}(X, \text{missed}) \quad (\in R^{B \times 31 \times 26}).$$

Let `posi` be the target position for example `i`. Extract

$$L = \text{logits}[:, 0, \dots, B-1, \text{pos}] \quad (\in R^{B \times 26}).$$

Add bias:

$$\text{bias_term} = \lambda_{\text{pos}} \log(\text{pos_bias}[:, \text{pos}]^\top + 1e^{-9}) \quad (\in R^{B \times 26}),$$

where `λpos = 0.2`. Then

$$L_{\text{biased}} = L + \text{bias_term}, \quad \text{loss} = \text{CrossEntropyLoss}(\text{label_smoothing} = 0.1)(L_{\text{biased}}, y_{\text{letter}}).$$

Backpropagation uses `autocast()` and `GradScaler()` for mixed precision.

6. Training Details

- Epochs: 18 (with early stopping if validation win-rate does not improve for 5 consecutive epochs).
- Optimizer: `AdamW(lr = 1 × 10-3, weight_decay = 1 × 10-2)`.
- Scheduler: `CosineAnnealingLR(Tmax = 50)`.
- Mixed precision via `torch.cuda.amp`.
- After each epoch, run full validation over 12 500 held-out words by simulating Hangman games (max 6 misses), compute win-rate, save best model.

Data Structures

- **CHAR_MAP:** $\{_: 27, \{ : 28, | : 29, a : 1, \dots, z : 26\}$. Index 0 is padding.
- `embedding_len` = 31. Maximum word length (29) + 2 boundary markers.
- `pos_bias` $\in R^{26 \times 31}$. Each column sums to 1 across 26 letters, representing $P(\text{letter} = c \mid \text{position} = p)$.
- $\vec{\epsilon} \in \{0, \dots, 29\}^{31}$ (input indices for `Embedding`).
- `missed_vec` $\in \{0, 1\}^{26}$ (binary mask of incorrect guesses).
- `revealed_mask` $\in \{0, 1\}^{31}$ (1 if that position is a revealed letter).
- `mask_pos` $\subset \{0, \dots, 30\}$ (indices of blanks).
- `target_pos` $\in \{0, \dots, 30\}$, `y_letter` $\in \{0, \dots, 25\}$.
- Batch-level tensors (B=256):

$$X \in Z^{256 \times 31}, \quad \text{missed} \in Z^{256 \times 26}, \quad \text{pos} \in Z^{256}, \quad y_letter \in Z^{256}.$$

- LSTM internal: input at each time step is 80-dim; output is (256, 31, 1024).
- Dropout applied to (256, 31, 1024) during training.
- Final FC: (256, 31, 1024) \rightarrow (256, 31, 26).

Code Snippet

```
def guess(self, word):
    clean = word.replace(' ', '')
    VOWELORDER = ['e', 'a', 'o', 'i', 'u']
    alpha = 0.6

    # Vowel-First Heuristic
    if (not any(v in clean for v in VOWELORDER)) and \
        (not any(g in self.guessed_letters for g in VOWELORDER)):
        for v in VOWELORDER:
            if v not in self.guessed_letters:
                self.guessed_letters.append(v)
        return v

    # Candidate Filtering via Regex
    pattern = ''.join(c if c != '_' else '.' for c in clean)
    possible_words = [
        w for w in self.full_dictionary
        if len(w) == len(clean) and re.match(f"^{pattern}$", w)
    ]
    candidate_letters = [c for c in string.ascii_lowercase
                        if c not in self.guessed_letters]

    # Frequency Scores
    if possible_words and candidate_letters:
```

```

letter_counts = {
    c: sum(1 for w in possible_words if c in w)
    for c in candidate_letters
}
max_count = max(letter_counts.values(), default=0)
freq_scores = {
    c: (letter_counts[c] / max_count) if max_count > 0 else 0.0
    for c in candidate_letters
}
else:
    freq_scores = {c: 0.0 for c in candidate_letters}

# Neural Model Prediction
padded = '{ ' + clean + ' | '
vec, missed_vec, revealed_mask, unrevealed_emb_positions = encode_input(
    padded, clean, self.guessed_letters, self.embedding_len
)
X = torch.tensor([vec], dtype=torch.long).to(self.device) # (1,50)
missed = torch.tensor([missed_vec], dtype=torch.long).to(self.device)
# (1,26)
with torch.no_grad():
    logits = self.model(X, missed)[0] # (50,26)
    probs = torch.softmax(logits, dim=-1)

    if unrevealed_emb_positions:
        probs_unrevealed = probs[unrevealed_emb_positions] # (num_blanks,26)
        total_probs = probs_unrevealed.sum(dim=0) # (26,)
        masked_total_probs = total_probs.clone()
        for g in self.guessed_letters:
            idx = ord(g) - 97
            if 0 <= idx < 26:
                masked_total_probs[idx] = 0
        model_scores = {
            chr(i + 97): masked_total_probs[i].item()
            for i in range(26)
        }
    else:
        model_scores = {c: 0.0 for c in candidate_letters}

# Combine Scores & Choose
combined_scores = {
    c: alpha * model_scores.get(c, 0.0)
    + (1 - alpha) * freq_scores.get(c, 0.0)
    for c in candidate_letters
}
if combined_scores:
    chosen_letter = max(combined_scores, key=combined_scores.get)
else:
    chosen_letter = 'a' # fallback

self.guessed_letters.append(chosen_letter)

```

```
return chosen_letter
```

Function Descriptions

Function	Description
<code>encode_input</code>	Given <code>masked_word</code> , <code>original_word</code> , <code>guessed_letters</code> , <code>embedding_len</code> , constructs: 1) $\vec{\in} \{0, \dots, 29\}^{\text{embedding_len}}$ (mapping <code> </code> , letters, “-”, padding), 2) <code>missed_vec</code> $\in \{0, 1\}^{26}$ (incorrect guesses), 3) <code>revealed_mask</code> $\in \{0, 1\}^{\text{embedding_len}}$ (1 if that position is a revealed letter), 4) <code>mask_pos</code> $\subset \{0, \dots, \text{embedding_len} - 1\}$ (indices of “-”).
<code>compute_pos_bias</code>	From list of words, frames each as <code>{w }</code> , tallies letter frequencies per position into <code>counts[26 × 31]</code> , normalizes columns to produce <code>pos_bias</code> $\in [0, 1]^{26 \times 31}$.
<code>HangmanDataset</code>	On each <code>__getitem__</code> , pads a word <code>{w }</code> , randomly reveals a fraction $\in \{0, 0.2, 0.5, 0.8\}$, masks some letters with “-”, calls <code>encode_input</code> , selects one masked position as target. Returns <code>(X, missed_vec, revealed_mask, y_letter, pos)</code> .
<code>HangmanLSTM</code>	Neural model: - <code>char_embed</code> : $(30 \rightarrow 64)$, - <code>missed_embed</code> : $(26 \rightarrow 16)$, - <code>lstm</code> : LSTM($80 \rightarrow 512 \times 2$, <code>layers</code> = 2), - <code>dropout</code> : $p = 0.5$, - <code>fc</code> : $(1024 \rightarrow 26)$. <code>forward(x, missed)</code> returns <code>(B, 31, 26)</code> logits.
<code>train_loop</code>	For each epoch: 1) Iterate over <code>dl_train</code> : compute <code>logits = model(X, missed) $\in (R^{B \times 31 \times 26})$</code> , extract <code>L = logits[:, pos] $\in R^{B \times 26}$</code> , add <code>$\lambda_{\text{pos}} \log(\text{pos_bias}[:, \text{pos}])$</code> , compute <code>CrossEntropyLoss(label_smoothing = 0.1)</code> , backprop with mixed precision. 2) Step scheduler, run validation on 12 500 words by simulating Hangman games (six misses), compute win-rate, save best model, early stop if no improvement for 5 epochs.
<code>guess</code>	Implements the hybrid inference pipeline: (a) Vowel-first for high information early; (b) Regex filter \rightarrow frequency scores; (c) LSTM \rightarrow neural scores; (d) Combined weighted score $\alpha = 0.6$. Returns the chosen letter.
<code>start_game</code>	Resets <code>guessed_letters = \emptyset</code> , calls API <code>"/new_game"</code> , loops until solved or 6 misses, calling <code>guess()</code> and sending <code>"/guess_letter"</code> each iteration, returns True if word is solved.

Table 1: Summary of key functions

Conclusion

This neural-statistical Hangman solver integrates:

- A **vowel-first heuristic** for rapid information gain.

- A **regex-based frequency filter** over dictionary candidates.
- A **bi-directional LSTM** (hidden=512, 2 layers, dropout=0.5) trained on synthetic masked examples.
- A **positional bias tensor** guiding the LSTM via biased loss.
- A **linear combination** of neural (=0.6) and frequency (0.4) scores at inference.

This solver achieves a **57.6 %** win-rate, outperforming other naive ML(CNN = 47 %, LSTM = 52 %) approaches. Future directions include dynamic scheduling, transformer-based encoders, and curriculum training from shorter to longer words.